

# Difference between `vm.dirty_ratio` and `vm.dirty_background_ratio`?

[Ask Question](#)

11



8

I'm currently experimenting with the kernel parameters found in `/proc/sys/vm`, especially `dirty_ratio` and `dirty_background_ratio`.

The kernel doc has the following explanations for both:

`dirty_background_ratio`

Contains, as a percentage of total available memory that contains free pages and reclaimable pages, the number of pages at which the background kernel flusher threads will start writing out dirty data.

and

`dirty_ratio`

Contains, as a percentage of total available memory that contains free pages and reclaimable pages, the number of pages at which a process which is generating disk writes will itself start writing out dirty data.

On my linux system `dirty_background_ratio` is 10 and `dirty_ratio` is 20. I understand that the difference is, who the dirty data writes. So if my used memory reaches 10% the kernel starts writing back and 20% should never be reached.

My question now is: Has the higher value of `dirty_background_ratio` and `dirty_ratio` any meaning or is it just a matter of "what is the lower value and who has it"?

[linux](#)[performance](#)[linux-kernel](#)

share improve this question

edited Mar 22 '15 at 11:34



Alexander Azarov

10.2k ● 1 ● 38 ● 43

asked Jan 12 '15 at 10:47



happyMOOyear

500 ● 1 ● 6 ● 13

add a comment

## 2 Answers

active oldest votes

▲  
21

Has the higher value of `dirty_background_ratio` and `dirty_ratio` any meaning or is it just a matter of "what is the lower value and who has it"?



In more simpler words:



**`vm.dirty_background_ratio`** is the percentage of system memory which when dirty then system can start writing data to the disks.

**`vm.dirty_ratio`** is percentage of system memory which when dirty, the process doing writes would block and write out dirty pages to the disks.

This tunable depends on what your system running, if you run large database its recommend to keep these values low to avoid I/O bottle-necks and when the system load increases.

ex:

```
vm.dirty_background_ratio=10
vm.dirty_ratio=15
```

In this example when the dirty pages exceed `vm.dirty_background_ratio=10` I/O starts, i.e they start getting flushed / written to the disk. When total number of dirty pages exceed `vm.dirty_ratio=15` all writes get blocked until some of the dirty pages get written to disk. You can think of the `vm.dirty_ratio=15` as the upper limit.

share improve this answer

edited Dec 30 '16 at 22:04



Mark Lakata

answered Jan 12 '15 at 12:34



askb



14.1k ● 4 ● 77 ● 93



4,165 ● 14 ● 31

- 3 So let me rephrase that, just to see if I understood correctly. If the `dirty_background_ratio` is reached, the kernel starts doing the writebacks in the background but applications can still write to the page cache without blocking. If `dirty_ratio` is reached, applications block on writing until `dirty_ratio` is no longer reached. Is that correct? – [happyMOOyear](#) Jan 12 '15 at 12:48
- 1 yes, I have also updated an use case if it helps. – [askb](#) Jan 12 '15 at 13:41
- I accepted your answer because it answers my question and helped me a lot. Thank you! Just one extra question: Is there some place where this behavior is documented? – [happyMOOyear](#) Jan 12 '15 at 13:59
- 2 "vm.dirty\_ratio is the value that represents the percentage of MemTotal that can consume dirty pages before all processes must write dirty buffers back to disk and when this value is reached all I/O is blocked for any new writes until dirty pages have been flushed." [sysxperits.com/home/announce/...](#) ... you can also verify this behaviour from the code. – [askb](#) Jan 12 '15 at 14:17
- Shouldn't "... vm.dirty\_ratio=10 as the upper limit" be "... vm.dirty\_ratio=15 as the upper limit" ? – [hbogert](#) Jun 22 '16 at 7:31

[show 1 more comment](#)



I have been intrigued by this very question and so experimented a bit on my Debian 7.10 system running Linux 3.2.0-4-amd64 using [sysbench](#) 0.4.12, modifying:

2

- `/proc/sys/vm/dirty_ratio`
- `/proc/sys/vm/dirty_background_ratio`

These settings are a way to delay writing to disk. They are useful as long as you have applications that write infrequently or in small chunks (e.g. web browser). If there is only one application on the system that is just generating data at a rate greater than the maximum supported by the disk then no settings matter. The writing will take as much time as it has to.

**Dirty Ratio** (DR) results in the process that caused the number of dirty pages to cross the threshold to block. **Dirty Background Ratio** (DBR) controls writing dirty pages in the background. So, if you have a low DBR, higher DR and all of your processes write in small chunks never in total crossing the supported write speed of the disk (e.g. 50 MB/s) then you will find a system that is pretty responsive. This is impressive when we bear in mind the fact that writing to RAM is usually 100 times faster (5 GB/s)! This is the importance of DBR.

Configuration parameters are useful when you are bothered about applications that write infrequently. You don't want a process writing a byte or reading a few KB to stall for 20 seconds because there is too much dirty data. This is the importance of not having a too high DR. It also ensures that some memory is available to cache recently used data.

[share](#) [improve this answer](#)

[edited Jun 5 '16 at 7:18](#)

[answered Jun 4 '16 at 13:55](#)



[pdp](#)

1,833 ● 1 ● 8 ● 14

[add a comment](#)

## Your Answer

**B** *I* [link](#) [quote](#) [code](#) [image](#) [list](#) [table](#) [undo](#) [redo](#)

Sign up or [log in](#)

 Sign up using Google

 Sign up using Facebook

 Sign up using Email and Password

## Post as a guest

**Name**

**Email**

Required, but never shown

Post Your Answer

By clicking "Post Your Answer", you acknowledge that you have read our updated [terms of service](#), [privacy policy](#) and [cookie policy](#), and that your continued use of the website is subject to these policies.

Not the answer you're looking for? Browse other questions tagged [linux](#) [performance](#) [linux-kernel](#) or [ask your own question](#).

asked 3 years, 10 months ago


viewed 26,993 times

active 1 year, 11 months ago

66 People Chatting

[SO](#) [Close](#) [Vote](#) [Reviewers](#)

1 hour ago - [Tetsuya Yamamoto](#)



[SOBotics](#)

1 hour ago - [FireAlarm](#)



## Related

710

[Efficiency of Java “Double Brace Initialization”?](#)

2649

[Improve INSERT-per-second performance of SQLite?](#)

2801

[What is the difference between call and apply?](#)

561

[Get difference between two lists](#)

6

[implementation of dirty\\_expire\\_centisecs](#)

92

[How can I find out the total physical memory \(RAM\) of my linux box suitable to be parsed by a shell script?](#)

1217

[Replacing a 32-bit loop counter with 64-bit introduces crazy performance deviations](#)

1

[File writes stall during heavy disk I/O on Linux box](#)



















101








[Writing programs to cope with I/O errors causing lost writes on Linux](#)

2

[Why the memory not fully used to cache data](#)

## Hot Network Questions

-  In the double-slit experiment, why is it never shown that particles may hit the space between or outside the slits?
-  Martial arts in zero or lower gravity
-  As a Kensei monk, can I use the Martial Arts feature to change the 1d4 extra damage from Kensei's Shot to my Martial Arts die?
-  No Q bar on flip-flop
-  Why don't marbles naturally arrange themselves like a crystal?
-  Can a cat who only lives in a (small) flat be happy?
-  How to prove ' $\exists xP(x)$ ' from ' $\neg\forall x(P(x)\rightarrow Q(x))$ '
-  Why is embassy spelled with E but ambassador with A?
-  Orange "aura" around people/objects in photo, usually present with older cameras?
-  What logical fallacy is "If you don't like it, move!"?
-  Squares for answer choice options and perfect alignment to mathematical answers
-  Is it possible to print the file size of a file?
-  How frustrating is my movie?
-  Could Sam have pierced Shelob's carapace?
-  Is there a relationship between the words "night" and "eight"?
-  What is the teacher's last name?
-  Copyright of name in movies
-  Nth subset of a set

-  [Proving Limit Rigorously](#)
-  [Could a name server resolve IP addresses dynamically based on some strategy?](#)
-  [Should I close port 80 forever and ever since the 2018 Google-indicated web-security initiatives?](#)
-  [Making sure if it is Cauchy](#)
-  [Delta carry-on policy](#)
-  [Creating an incompetent antagonist](#)
-  [question feed](#)



## STACK OVERFLOW

[Questions](#)  
[Jobs](#)  
[Developer Jobs Directory](#)  
[Salary Calculator](#)  
[Help](#)  
[Mobile](#)  
[Disable Responsiveness](#)

## PRODUCTS

[Teams](#)  
[Talent](#)  
[Engagement](#)  
[Enterprise](#)

## COMPANY

[About](#)  
[Press](#)  
[Work Here](#)  
[Legal](#)  
[Privacy Policy](#)  
[Contact Us](#)

## STACK EXCHANGE NETWORK

[Technology >](#)  
[Life / Arts >](#)  
[Culture / Recreation >](#)  
[Science >](#)  
[Other >](#)

[Blog](#)
[Facebook](#)
[Twitter](#)
[LinkedIn](#)

site design / logo © 2018 Stack Exchange Inc; user  
 contributions licensed under cc by-sa 3.0 with  
 attribution required. rev 2018.11.22.32249